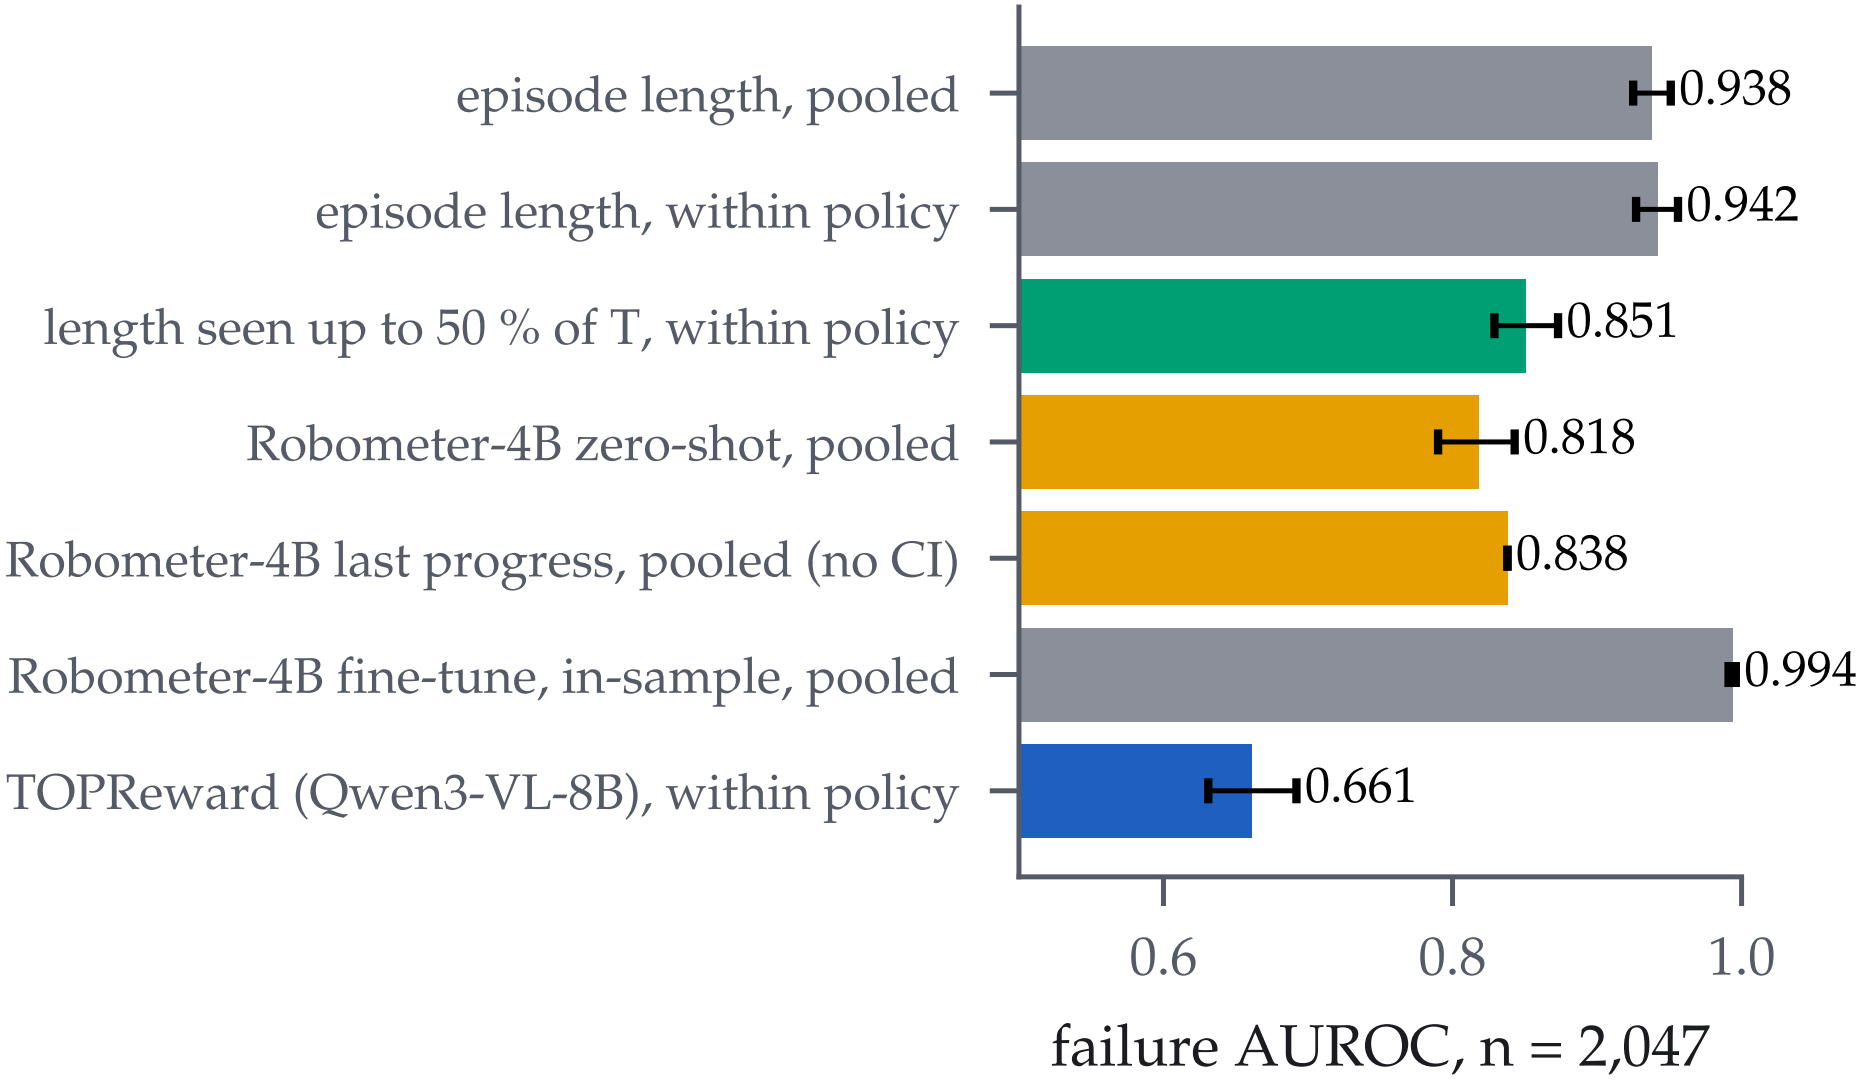
episode length, pooled
episode length, within policy
length seen up to 50 % of T, within policy
Robometer-4B zero-shot, pooled
Robometer-4B last progress, pooled (no CI)
Robometer-4B fine-tune, in-sample, pooled
TOPReward (Qwen3-VL-8B), within policy
0.938
0.942
0.851
0.818
0.838
0.994
0.661
0.6
0.8
1.0
failure AUROC, n = 2,047